

ML/DL for Everyone Season2

with  TensorFlow

Lab 07-1 Application & Tips

Code: <https://github.com/deeplearningzerotoall/TensorFlow>

Slides: <http://bit.ly/2LQMKVv>

Lecturer: SuSang Kim (healess@kaist.ac.kr)

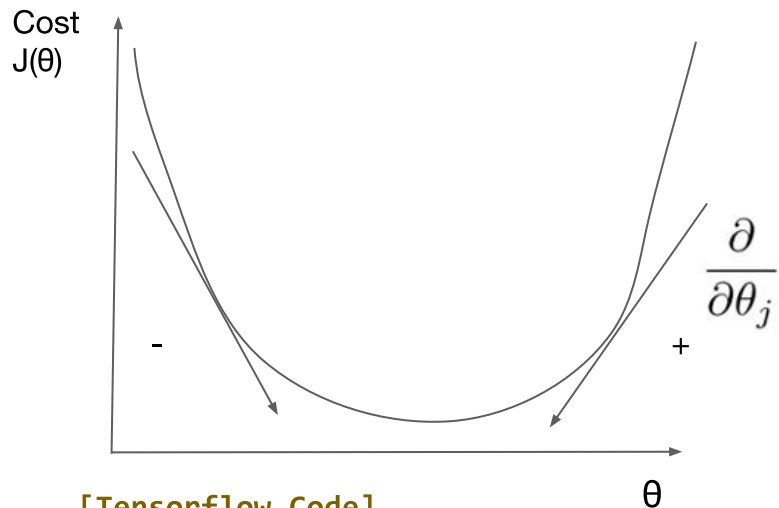


Application & Tips

- Learning rate
 - Gradient
 - Good and Bad Learning rate
 - Annealing the learning rate (Decay)
- Data preprocessing
 - Standardization / Normalization
 - Noisy Data
- Overfitting
 - Regularization
 - L2 Norm

Learning rate

Gradient



$$\text{Repeat } \{ \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \}$$


Learning rate is a hyper-parameter that controls how much we are adjusting the weights with respect the loss gradient

[Tensorflow Code]

```
def grad(hypothesis, labels):  
    with tf.GradientTape() as tape:  
        loss_value = loss_fn(hypothesis, labels)  
    return tape.gradient(loss_value, [W,b])  
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)  
optimizer.apply_gradients(grads_and_vars=zip(grads,[W,b]))
```


Learning rate

Good and Bad



[Train Log]

Iter: 0, Loss: 6.0257, Learning Rate: 0.1000
Iter: 1000, Loss: 0.3723, Learning Rate: 0.0960
Iter: 2000, Loss: 0.2779, Learning Rate: 0.0922
Iter: 3000, Loss: 0.2293, Learning Rate: 0.0885
Iter: 4000, Loss: 0.1977, Learning Rate: 0.0849
Iter: 5000, Loss: 0.1750, Learning Rate: 0.0815



High Learning Rate is Overshooting

Normal Learning Rate is **0.01**



3e-4 is the best learning rate for Adam,
hands down (andrey karpathy)

[Tensorflow Code]

```
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
```

Learning rate

Annealing the learning rate



Log

```
Iter: 0, Loss: 1.7346, Learning Rate: 0.10000000
Iter: 10, Loss: 0.0745, Learning Rate: 0.10000000
Iter: 20, Loss: 0.0438, Learning Rate: 0.10000000
Iter: 30, Loss: 0.0273, Learning Rate: 0.10000000
Iter: 40, Loss: 0.0181, Learning Rate: 0.10000000
Iter: 50, Loss: 0.0128, Learning Rate: 0.09600000
Iter: 60, Loss: 0.0099, Learning Rate: 0.09600000
Iter: 70, Loss: 0.0080, Learning Rate: 0.09600000
Iter: 80, Loss: 0.0068, Learning Rate: 0.09600000
Iter: 90, Loss: 0.0060, Learning Rate: 0.09600000
Iter: 100, Loss: 0.0054, Learning Rate: 0.09216000
```




[Tensorflow Code]

```
learning_rate = tf.train.exponential_decay(starter_learning_rate,
                                           global_step, 1000, 0.96, staircase=True)
# tf.train.exponential_decay / tf.train.inverse_time_decay
# tf.train.natural_exp_decay / tf.train.piecewise_constant
# tf.train.polynomial_decay
```

```
def exponential_decay(epoch):
    starter_rate = 0.01
    k = 0.96
    exp_rate = starter_rate * exp(-k*t)
    return exp_rate
```

Data preprocessing

Feature Scaling



Standardization
(Mean Distance)

$$x_{new} = \frac{x - \mu}{\sigma}$$

[Python Code(numpy)]

```
Standardization = (data - np.mean(data)) / sqrt(np.sum(  
(data - np.mean(data))^2 ) / np.count(data))
```


Normalization
(0~1)

$$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$$


```
Normalization = (data - np.min(data, 0)) / (np.max(data, 0)  
- np.min(data, 0))
```

Data preprocessing

Noisy Data




Numeric



Will you order a pizza..??

NLP

You order pizza



Face
Image

