


ML/DL for Everyone Season2

with  TensorFlow

Lab 11-5 Best CNN with MNIST Dataset



Code: <https://github.com/deeplearningzerotoall/TensorFlow>

Slides: <http://bit.ly/2LQMKVv>

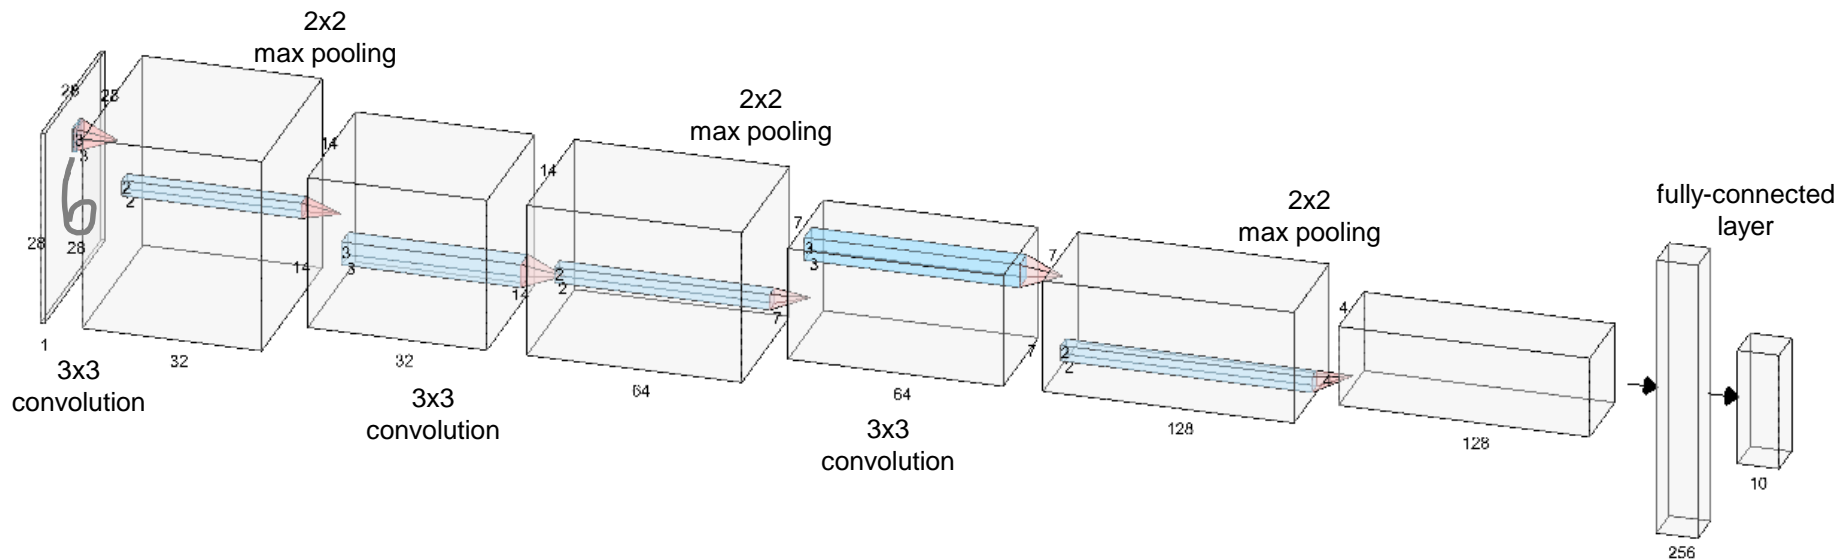
Lecturer: 이진원/JinWon Lee(wlee.ml25@gmail.com)



NN Implementation Flow in TensorFlow

1. Set hyper parameters – learning rate, training epochs, batch size, etc.
2. Data Augmentation – rotate & shift
3. Make a data pipelining – use `tf.data`
4. Build a neural network model – use `tf.keras`
5. Define a loss function – cross entropy
6. Calculate a gradient – use `tf.GradientTape`
7. Select an optimizer – Adam optimizer
8. Define a metric for model's performance – accuracy
9. (optional) Make a checkpoint for saving
10. Train and Validate a neural network model

CNN with MNIST Data



0. Import Libraries

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.utils import to_categorical
import numpy as np
import matplotlib.pyplot as plt
import os
from scipy import ndimage
```

1. Set Hyper Parameters

```
learning_rate = 0.001  
training_epochs = 15  
batch_size = 100
```

```
cur_dir = os.getcwd()  
ckpt_dir_name = 'checkpoints'  
model_dir_name = 'minst_cnn_best'
```

```
checkpoint_dir = os.path.join(cur_dir, ckpt_dir_name, model_dir_name)  
os.makedirs(checkpoint_dir, exist_ok=True)
```

```
checkpoint_prefix = os.path.join(checkpoint_dir, model_dir_name)
```

2. Data Augmentation

```
from scipy import ndimage
```

```
def data_augmentation(images, labels):  
    aug_images = []; aug_labels = []  
  
    for x, y in zip(images, labels):  
        aug_images.append(x)  
        aug_labels.append(y)  
  
        bg_value = np.median(x)  
        for _ in range(4):  
            angle = np.random.randint(-15, 15, 1)  
            rot_img = ndimage.rotate(x, angle, reshape=False, cval=bg_value)  
  
            shift = np.random.randint(-2, 2, 2)  
            shift_img = ndimage.shift(rot_img, shift, cval=bg_value)  
  
            aug_images.append(shift_img)  
            aug_labels.append(y)  
    aug_images = np.array(aug_images)  
    aug_labels = np.array(aug_labels)  
    return aug_images, aug_labels
```

3. Make a Data Pipelining

```
mnist = keras.datasets.mnist
```

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
train_images, train_labels = data_augmentation(train_images, train_labels)
```

```
train_images = train_images.astype(np.float32) / 255.
```

```
test_images = test_images.astype(np.float32) / 255.
```

```
train_images = np.expand_dims(train_images, axis=-1)
```

```
test_images = np.expand_dims(test_images, axis=-1)
```

```
train_labels = to_categorical(train_labels, 10)
```

```
test_labels = to_categorical(test_labels, 10)
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_images,
                                                    train_labels)).shuffle(buffer_size=500000).batch(batch_size)
```

[illegible]

4. Build a Neural Network Model

```
class ConvBNRelu(tf.keras.Model):  
    def __init__(self, filters, kernel_size=3, strides=1, padding='SAME'):  
        super(ConvBNRelu, self).__init__()  
        self.conv = keras.layers.Conv2D(filters=filters, kernel_size=kernel_size,  
                                          strides=strides, padding=padding,  
                                          kernel_initializer='glorot_normal')  
        self.batchnorm = tf.keras.layers.BatchNormalization()  
  
    def call(self, inputs, training=False):  
        layer = self.conv(inputs)  
        layer = self.batchnorm(layer)  
        layer = tf.nn.relu(layer)  
        return layer
```


4. Build a Neural Network Model

```
class DenseBNRelu(tf.keras.Model):  
    def __init__(self, units):  
        super(DenseBNRelu, self).__init__()  
        self.dense = keras.layers.Dense(units=units,  
                                          kernel_initializer='glorot_normal')  
        self.batchnorm = tf.keras.layers.BatchNormalization()  
  
    def call(self, inputs, training=False):  
        layer = self.dense(inputs)  
        layer = self.batchnorm(layer)  
        layer = tf.nn.relu(layer)  
        return layer
```

4. Build a Neural Network Model

[illegible]

4. Build a Neural Network Model

```
def call(self, inputs, training=False):  
    net = self.conv1(inputs)  
    net = self.pool1(net)  
    net = self.conv2(net)  
    net = self.pool2(net)  
    net = self.conv3(net)  
    net = self.pool3(net)  
    net = self.pool3_flat(net)  
    net = self.dense4(net)  
    net = self.drop4(net)  
    net = self.dense5(net)  
    return net
```

```
models = []  
num_models = 5  
for m in range(num_models):  
    models.append(MNISTModel())
```

5. Define a Loss Function

6. Calculate a Gradient

```
def loss_fn(model, images, labels):  
    logits = model(images, training=True)  
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(  
        logits=logits, labels=labels))  
    return loss
```

```
def grad(model, images, labels):  
    with tf.GradientTape() as tape:  
        loss = loss_fn(model, images, labels)  
    return tape.gradient(loss, model.variables)
```

7. Select an Optimizer

8. Define a Metric for Model's Performance

9. Make a Checkpoint for Saving

```
global_step = tf.train.get_or_create_global_step()
lr_decay = tf.train.exponential_decay(learning_rate, global_step,
                                     train_images.shape[0]/batch_size*num_models*5,
                                     0.5, staircase=True)
optimizer = tf.train.AdamOptimizer(learning_rate=lr_decay)
```

```
def evaluate(models, images, labels):
    predictions = tf.zeros_like(labels)
    for model in models:
        logits = model(images, training=False)
        predictions += logits
    correct_prediction = tf.equal(tf.argmax(predictions, 1), tf.argmax(labels, 1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    return accuracy
```

```
checkpoints = []
for m in range(num_models):
    checkpoints.append(tf.train.Checkpoint(cnn=models[m]))
```

10. Train and Validate a Neural Network Model

```
for epoch in range(training_epochs):
    avg_loss = 0.; avg_train_acc = 0.; avg_test_acc = 0.
    train_step = 0; test_step = 0

    for images, labels in train_dataset:
        for model in models:
            grads = grad(model, images, labels)
            optimizer.apply_gradients(zip(grads, model.variables))
            loss = loss_fn(model, images, labels)
            avg_loss += loss / num_models
        acc = evaluate(models, images, labels)
        avg_train_acc += acc
        train_step += 1
    avg_loss = avg_loss / train_step
    avg_train_acc = avg_train_acc / train_step
```

10. Train and Validate a Neural Network Model

```
for images, labels in test_dataset:
    acc = evaluate(models, images, labels)
    avg_test_acc += acc
    test_step += 1
avg_test_acc = avg_test_acc / test_step

print('Epoch:', '{}'.format(epoch + 1), 'loss =', '{:.8f}'.format(avg_loss),
      'train accuracy = ', '{:.4f}'.format(avg_train_acc),
      'test accuracy = ', '{:.4f}'.format(avg_test_acc))

for idx, checkpoint in enumerate(checkpoints):
    checkpoint.save(file_prefix=checkpoint_prefix+'-{}'.format(idx))
```

Accuracy : 99.68%

What's Next?

- Recurrent Neural Network (RNN)