

ML/DL for Everyone Season2

with  TensorFlow

Lab 12-5: Seq2Seq

Code: <https://github.com/deeplearningzerotoall/TensorFlow>

Slides: https://drive.google.com/drive/folders/1twBsdLkl2P15J0DqYs77_E_EVKt7Ghav

Lecturer: sungjin7127@gmail.com




Lab12-5: Seq2Seq


- Seq2Seq Overview
 - Chatbot Example
 - Encoder - Decoder
- Data Pipeline
- Encoder-Decoder
- Loss & Optimizer
- Train
- Prediction
- What's Next?

Seq2Seq Overview

What is the difference between general RNN model and Seq2Seq model?



RNN



Seq2Seq

Example : Chatbot

I broke up yesterday

Sorry to hear that.


----- After some conversation... -----

Today's perfect weather
makes me much sad

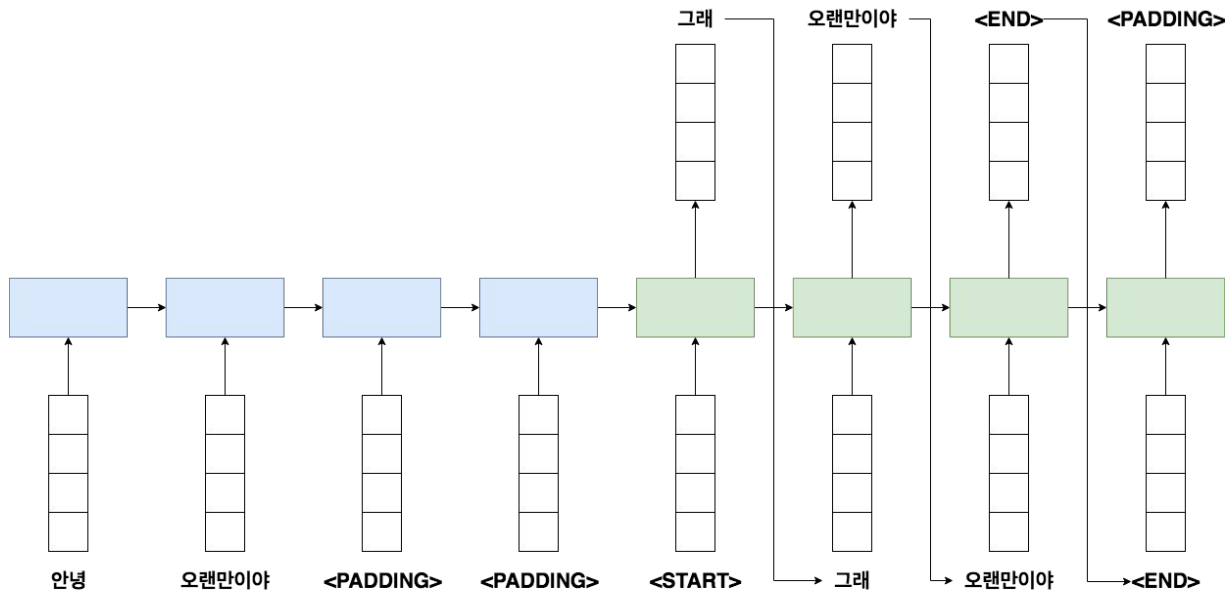


Today's perfect weather makes me much **sad**

Encoder-Decoder



Encoder-Decoder



Data Pipeline: Dataset

```
sources = [['I', 'feel', 'hungry'],  
           ['tensorflow', 'is', 'very', 'difficult'],  
           ['tensorflow', 'is', 'a', 'framework', 'for', 'deep', 'learning'],  
           ['tensorflow', 'is', 'very', 'fast', 'changing']]  
targets = [['나는', '배가', '고프다'],  
           ['텐서플로우는', '매우', '어렵다'],  
           ['텐서플로우는', '딥러닝을', '위한', '프레임워크이다'],  
           ['텐서플로우는', '매우', '빠르게', '변화한다']]
```

Data Pipeline: Vocab Dict

```
# vocabulary for sources
```

```
s_vocab = list(set(sum(sources, [])))
```

```
s_vocab.sort()
```

```
s_vocab = ['<pad>'] + s_vocab
```

```
source2idx = {word : idx for idx, word in enumerate(s_vocab)}
```

```
idx2source = {idx : word for idx, word in enumerate(s_vocab)}
```

```
pprint(source2idx)
```

```
# vocabulary for targets
```

```
t_vocab = list(set(sum(targets, [])))
```

```
t_vocab.sort()
```

```
t_vocab = ['<pad>', '<bos>', '<eos>'] + t_vocab
```

```
target2idx = {word : idx for idx, word in enumerate(t_vocab)}
```

```
idx2target = {idx : word for idx, word in enumerate(t_vocab)}
```

```
pprint(target2idx)
```

```
{'<pad>': 0,  
'l': 1,  
'a': 2,  
'changing': 3,  
'deep': 4,  
'difficult': 5,  
'fast': 6,
```

```
'feel': 7,  
'for': 8,  
'framework': 9,  
'hungry': 10,  
'is': 11,  
'learning': 12,  
'tensorflow': 13,  
'very': 14}
```

```
{'<bos>': 1,  
'<eos>': 2,  
'<pad>': 0,  
'고프다': 3,  
'나는': 4,  
'딥러닝을': 5,
```

```
'매우': 6,  
'배가': 7,  
'변화한다': 8,  
'빠르게': 9,  
'어렵다': 10,  
'위한': 11,  
'텐서플로우는': 12,  
'프레임워크이다': 13}
```


Data Pipeline: Preprocess (1/3)

```
def preprocess(sequences, max_len, dic, mode = 'source'):
    assert mode in ['source', 'target'], 'source와 target 중에 선택해주세요.'

    if mode == 'source':
        # preprocessing for source (encoder)
        s_input = list(map(lambda sentence : [dic.get(token) for token in sentence], sequences))
        s_len = list(map(lambda sentence : len(sentence), s_input))
        s_input = pad_sequences(sequences = s_input, maxlen = max_len, padding = 'post',
                               truncating = 'post')

    return s_len, s_input
```

⋮

Data Pipeline: Preprocess (2/3)

```
elif mode == 'target':  
    # preprocessing for target (decoder)  
    # input  
    t_input = list(map(lambda sentence : ['<bos>'] + sentence + ['<eos>'], sequences))  
    t_input = list(map(lambda sentence : [dic.get(token) for token in sentence], t_input))  
    t_len = list(map(lambda sentence : len(sentence), t_input))  
    t_input = pad_sequences(sequences = t_input, maxlen = max_len, padding = 'post',  
                           truncating = 'post')  
  
    # output  
    t_output = list(map(lambda sentence : sentence + ['<eos>'], sequences))  
    t_output = list(map(lambda sentence : [dic.get(token) for token in sentence], t_output))  
    t_output = pad_sequences(sequences = t_output, maxlen = max_len, padding = 'post',  
                           truncating = 'post')  
  
    return t_len, t_input, t_output
```

Data Pipeline: Preprocess (3/3)

preprocessing for source

s_max_len = 10

```
s_len, s_input = preprocess(sequences = sources,  
                             max_len = s_max_len, dic = source2idx, mode = 'source')
```

preprocessing for target

t_max_len = 12

```
t_len, t_input, t_output = preprocess(sequences = targets,  
                                       max_len = t_max_len, dic = target2idx, mode = 'target')
```

S_len: [3, 4, 7, 5]

S_input:

```
[[ 1 7 10 0 0 0 0 0 0 0]  
[13 11 14 5 0 0 0 0 0 0]  
[13 11 2 9 8 4 12 0 0 0]  
[13 11 14 6 3 0 0 0 0 0]]
```

t_len: [5, 5, 6, 6]

t_input:


```
[[ 1 4 7 3 2 0 0 0 0 0 0 0]  
[ 1 12 6 10 2 0 0 0 0 0 0 0]  
[ 1 12 5 11 13 2 0 0 0 0 0 0]  
[ 1 12 6 9 8 2 0 0 0 0 0 0]]
```

t_output

```
[[ 4 7 3 2 0 0 0 0 0 0 0 0]  
[12 6 10 2 0 0 0 0 0 0 0 0]  
[12 5 11 13 2 0 0 0 0 0 0 0]  
[12 6 9 8 2 0 0 0 0 0 0 0]]
```

Data Pipeline: tf.data

```
# input  
data = tf.data.Dataset.from_tensor_slices((s_len, s_input, t_len, t_input, t_output))  
data = data.shuffle(buffer_size = buffer_size)  
data = data.batch(batch_size = batch_size)  
# s_mb_len, s_mb_input, t_mb_len, t_mb_input, t_mb_output = iterator.get_next()
```




Encoder-Decoder

```
def gru(units):  
    # If you have a GPU, we recommend using CuDNNGRU(provides a 3x speedup than GRU)  
    # the code automatically does that.  
    if tf.test.is_gpu_available():  
        return tf.keras.layers.CuDNNGRU(units,  
                                           return_sequences=True,  
                                           return_state=True,  
                                           recurrent_initializer='glorot_uniform')  
    else:  
        return tf.keras.layers.GRU(units,  
                                     return_sequences=True,  
                                     return_state=True,  
                                     recurrent_activation='sigmoid',  
                                     recurrent_initializer='glorot_uniform')
```

Encoder-Decoder: Encoder


```
class Encoder(tf.keras.Model):  
    def __init__(self, vocab_size, embedding_dim, enc_units, batch_sz):  
        super(Encoder, self).__init__()  
        self.batch_sz = batch_sz  
        self.enc_units = enc_units  
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)  
        self.gru = gru(self.enc_units)  
  
    def call(self, x, hidden):  
        x = self.embedding(x)  
        output, state = self.gru(x, initial_state = hidden)  
  
        return output, state  
  
    def initialize_hidden_state(self):  
        return tf.zeros((self.batch_sz, self.enc_units))
```



Encoder-Decoder: Decoder(1/2)

```
class Decoder(tf.keras.Model):  
    def __init__(self, vocab_size, embedding_dim, dec_units, batch_sz):  
        super(Decoder, self).__init__()  
        self.batch_sz = batch_sz  
        self.dec_units = dec_units  
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)  
        self.gru = gru(self.dec_units)  
        self.fc = tf.keras.layers.Dense(vocab_size)
```


⋮



Encoder-Decoder: : Decoder(2/2)

```
def call(self, x, hidden, enc_output):  
  
    x = self.embedding(x)  
    output, state = self.gru(x, initial_state = hidden)  
    # output shape == (batch_size * 1, hidden_size)  
    output = tf.reshape(output, (-1, output.shape[2]))  
    # output shape == (batch_size * 1, vocab)  
    x = self.fc(output)  
  
    return x, state  
  
def initialize_hidden_state(self):  
    return tf.zeros((self.batch_sz, self.dec_units))
```

```
encoder = Encoder(len(source2idx), embedding_dim, units, batch_size)  
decoder = Decoder(len(target2idx), embedding_dim, units, batch_size)
```



Loss & Optimizer

```
def loss_function(real, pred):  
    mask = 1 - np.equal(real, 0)  
    loss_ = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=real, logits=pred) * mask  
  
    return tf.reduce_mean(loss_)  
  
# creating optimizer  
opt = tf.train.AdamOptimizer()  
  
# creating check point (Object-based saving)  
checkpoint_dir = './data_out/training_checkpoints'  
checkpoint_prefix = os.path.join(checkpoint_dir, 'ckpt')  
checkpoint = tf.train.Checkpoint(optimizer=optimizer,  
                                encoder=encoder,  
                                decoder=decoder)  
  
# create writer for tensorboard  
summary_writer = tf.contrib.summary.create_file_writer(logdir=checkpoint_dir)
```

Train (1/4)

EPOCHS = 100

```
for epoch in range(EPOCHS):
    hidden = encoder.initialize_hidden_state()
    total_loss = 0
    for i, (s_len, s_input, t_len, t_input, t_output) in enumerate(data):
        loss = 0
        with tf.GradientTape() as tape:
            enc_output, enc_hidden = encoder(s_input, hidden)
            dec_hidden = enc_hidden
            dec_input = tf.expand_dims([target2idx['<bos>']], *batch_size, 1)
            #Teacher Forcing: feeding the target as the next input
            for t in range(1, t_input.shape[1]):
                predictions, dec_hidden = decoder(dec_input, dec_hidden, enc_output)
                loss += loss_function(t_input[:, t], predictions)
                dec_input = tf.expand_dims(t_input[:, t], 1) #using teacher forcing
```

⋮

Train (2/4)

I feel hungry

W/O Teacher Forcing

X

Y predict

1. [bos]

a

2. [bos], a

?

Teacher Forcing

1. [bos]

?

2. [bos], I

?

3. [bos], I feel

?

Train (3/4)

```
batch_loss = (loss / int(t_input.shape[1]))
total_loss += batch_loss
variables = encoder.variables + decoder.variables
gradient = tape.gradient(loss, variables)
optimizer.apply_gradients(zip(gradient, variables))

if epoch % 10 == 0:
    #save model every 10 epoch
    print('Epoch {} Loss {:.4f} Batch Loss {:.4f}'.format(epoch,
                                                            total_loss / n_batch,
                                                            batch_loss.numpy()))
    checkpoint.save(file_prefix = checkpoint_prefix)
```

Train (4/4)

Epoch 0 Loss 0.0307 Batch Loss 0.7687
Epoch 10 Loss 0.0297 Batch Loss 0.7414
Epoch 20 Loss 0.0267 Batch Loss 0.6676
Epoch 30 Loss 0.0237 Batch Loss 0.5925
Epoch 40 Loss 0.0159 Batch Loss 0.3964
Epoch 50 Loss 0.0131 Batch Loss 0.3271
Epoch 60 Loss 0.0100 Batch Loss 0.2498
Epoch 70 Loss 0.0075 Batch Loss 0.1874
Epoch 80 Loss 0.0051 Batch Loss 0.1283
Epoch 90 Loss 0.0031 Batch Loss 0.0763

⋮

Prediction (1/3)

```
def prediction(sentence, encoder, decoder, inp_lang, targ_lang, max_length_inp, max_length_targ):
```

```
    inputs = [inp_lang[i] for i in sentence.split(' ')]
    inputs = tf.keras.preprocessing.sequence.pad_sequences([inputs], maxlen=max_length_inp,
padding='post')
    inputs = tf.convert_to_tensor(inputs)
```

```
    result = ''
```

```
    hidden = [tf.zeros((1, units))]
    enc_out, enc_hidden = encoder(inputs, hidden)
```

```
    dec_hidden = enc_hidden
    dec_input = tf.expand_dims([targ_lang['<bos>']], 0)
```

```
    :
```

Prediction (2/3)

```
for t in range(max_length_targ):
    predictions, dec_hidden = decoder(dec_input, dec_hidden, enc_out)

    predicted_id = tf.argmax(predictions[0]).numpy()

    result += idx2target[predicted_id] + ' '

    if idx2target.get(predicted_id) == '<eos>':
        return result, sentence

    # the predicted ID is fed back into the model
    dec_input = tf.expand_dims([predicted_id], 0)

return result, sentence
```

Prediction (3/3)

```
sentence = 'tensorflow is a framework for deep learning'  
sentence = 'I feel hungry'
```

```
result, output_sentence = prediction(sentence, encoder, decoder, source2idx, target2idx, s_max_len,  
t_max_len)
```

Result: tensorflow is a framework for deep learning

Output Sentence: 텐서플로우는 딥러닝을 위한 프레임워크이다 <eos>

Result: I feel hungry

Output Sentence: 나는 배가 고프다 <eos>

What's Next?

- Seq2Seq Attention

Reference

- Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation: <https://arxiv.org/abs/1406.1078>
- 텐서플로우와 머신러닝으로 시작하는 자연어처리 (Wikibooks): <http://wikibook.co.kr/nlp/>